

QA-MATRIX: AUTONOMOUS QA TESTER & AUTOMATION RAG ENGINE

End-to-End System Architecture, Engineering Specifications & Workflow Blueprint

Author: Pursharth Singh

Core Stack: FastAPI · Gemini 3.6 Flash · Pinecone RAG

Version: 1.0.0 (Production Ready)


1. Executive Summary & Problem Resolution

QA-Matrix is an autonomous Quality Engineering intelligence platform designed to eliminate the latency between product specifications and test validation. By synthesizing 768-dimensional Vector RAG with Google Gemini 3.6 Flash reasoning, it autonomously ingests Product Requirement Documents (PRDs) and OpenAPI specifications to generate exhaustive test matrices, production-grade Playwright/Cypress automation suites, and defect root-cause diagnoses.

Key Architectural Objective: Ground 100% of generated test assertions, edge-case matrices, and code locators strictly against indexed requirement documents, preventing hallucination while accelerating test authoring velocity by over 90%.

2. End-to-End System Architecture Blueprint

The system is partitioned into six decoupled, high-performance layers:



3. Why RAG is Used & Its Architectural Benefits

In modern QA engineering, feeding monolithic 50+ page PRDs or large OpenAPI schemas directly into LLM prompts introduces catastrophic failure modes: high token expenses, sluggish latency (>15s), context dilution ('needle-in-a-haystack' degradation), and fabricated test cases (hallucinations). **QA-Matrix implements Retrieval-Augmented Generation (RAG)** to decouple requirement storage from real-time reasoning.

- **The Core Concept in Simple Terms (The 'Open-Book Exam' Principle):**
 - **Without RAG (Closed Book):** The AI tries to memorize or guess product rules, leading to fictitious endpoints and invalid test criteria.
 - **With RAG (Open Book):** When asked to generate a test or diagnose a bug, the AI instantly flips to the exact 2–3 requirement paragraphs in your PRD and writes code strictly conforming to those exact rules.

Core Technical Benefits of RAG in QA-Matrix:

- **100% Grounded Test Assertions:** Tests reflect real schema types, required parameters, and explicit acceptance criteria rather than generic estimates.
- **Pinpoint Root Cause Analysis (RCA):** On runtime failure, RAG matches the stack trace against exact PRD clauses, isolating the specific requirement violated and generating Jira bug tickets.
- **Sub-Second Retrieval & 90%+ Cost Reduction:** Embedding similarity filters out 98% of irrelevant documentation, passing only Top-4 context chunks (relevance ≥ 0.75) into the prompt.
- **Instant Zero-Retraining Adaptation:** Updating a PRD or API schema requires only re-indexing vectors into Pinecone/Local store; test generators update immediately without fine-tuning.
- **High-Fidelity Code Locators:** Injects exact field names and UI labels into Page Object Models (``getByRole``, ``getByTestId``), producing robust, flakiness-free automation scripts.

Evaluation Dimension	Standard LLM Prompting	QA-Matrix Vector RAG
Test Accuracy & Grounding	Hallucinates missing fields; invents routes.	100% strict ground-truth tied to PRD chunks.
Document Size Limit	Restricted by context window / token budget.	Unlimited; handles hundreds of PDFs & specs.
Token Efficiency & Cost	Re-transmits 50,000+ tokens on every click.	Transmits only ~1,200 tokens (Top-4 chunks).
Defect RCA Traceability	Generic troubleshooting advice.	Direct PRD clause citation + reproduction patch.
Spec Updates & Maintenance	Manual copy-pasting of updated documents.	Instant vector re-indexing via <code>/api/ingest-spec</code> .

4. Subsystem Technical Specifications

A. Document Intake & 768-Dimensional Vectorization Engine

- **Multi-Modal Parsing:** Ingests PDF specifications (via PyPDF stream parser), raw OpenAPI/Swagger JSON/YAML, or plain text user stories.
- **Semantic Chunking:** Partitions text into dynamic windows of 800 words with 100-word overlaps to preserve cross-sentence acceptance criteria.
- **Embedding Model:** Generates 768-dimensional vector representations via Google gemini-embedding-001.
- **Dual Storage:** Upserts vectors into Pinecone Serverless (namespace qa-specs-v1) with real-time replication to a high-speed local cosine-similarity vector store.

B. RAG Semantic Retrieval & Grounding Layer

- Calculates normalized dot-product cosine similarity: $Sim(u, v) = (u \cdot v) / (||u|| \cdot ||v||)$.
- Retrieves top-4 contextual chunks with semantic relevance scores and injects them directly into the Gemini prompt template, ensuring complete traceability to the source PRD.

C. Gemini 3.6 Flash Intelligence Core

- Executes reasoning at low temperature (0.15) with strict JSON schema enforcement.
- Evaluates failure modes, boundary lengths, authentication bypass vectors, and rate-limiting thresholds.

5. Visual End-to-End Workflow Lifecycle

The flowchart below illustrates how an incoming PRD or runtime failure traverses the QA-Matrix intelligence pipeline:



Detailed Workflow Steps:

- Step 1 — Document Intake & Indexing:** Uploads PRD PDF or schema. Ingestion engine computes embeddings and indexes vectors into RAG knowledge base.
- Step 2 — Context Retrieval & Matrix Gen:** Vector search finds matching spec chunks. Gemini Flash synthesizes Functional, Boundary, and Security test cards with P0-P3 priority ratings.
- Step 3 — Automation Code Synthesis:** Test cases compile into production Playwright (TS/Py), Cypress, PyTest, or Postman suites featuring clean Page Object Model design.
- Step 4 — Failure RCA & Live Verification:** Stack traces from failed runs are analyzed against the indexed PRD to isolate defect root causes and generate ready-to-publish Jira bug tickets.

6. Complete REST API Specifications

Endpoint	Method	Input Parameters	Function & Output
/api/ingest-spec	POST	spec_title, category, file / content_text	Chunks & vectorizes PRD documents into Pinecone & local store.
/api/generate-test-matrix	POST	feature_name, test_scope, requirements	Queries RAG store and synthesizes full P0-P3 test plan JSON.
/api/generate-automation-scripts	POST	feature_name, framework, test_cases	Synthesizes Playwright / Cypress / PyTest automation code.
/api/analyze-bug-log	POST	error_log, test_context, expected_behavior	Performs Root Cause Analysis and generates Jira defect ticket.
/api/execute-api-test	POST	url, method, headers, expected_status	Executes live asynchronous HTTP request and assertion checks.
/api/history	GET	None	Retrieves persistent history of test suites and bug reports.

7. CI/CD Orchestration & Deployment Guide

QA-Matrix includes an enterprise workflow export (qa-automated-tester-rag.json) designed for seamless integration with GitHub/GitLab webhooks and Slack notifications. When a pull request is opened, the webhook automatically queries the QA RAG engine, generates test cases, compiles Playwright suites, and broadcasts an executive quality report to Slack.

Deployment Commands & Endpoints:

- **Start Server:** `cd /Users/pursharthsingh/Desktop/Satyarth/Q&A;_tool && python3 app.py`
- **Web Dashboard:** `http://127.0.0.1:8001`
- **Zero Setup:** Pre-configured with Google Gemini 3.6 Flash & Pinecone credentials with robust local vector failover.